

# Docker Overview

jnagal@

Containerizing everything @ Google

Containers at scale.

Resource Isolation.

[Imctfy](#)

[libcontainer](#)



# Docker : What & Why

Machine or Application containers

Build Once, Configure Once.

Deploy Everything\*

Everywhere\*

Reliably & Consistently

Efficiently

Cheaply

# Docker Features

Image  
Management

Resource  
Isolation

File system  
Isolation

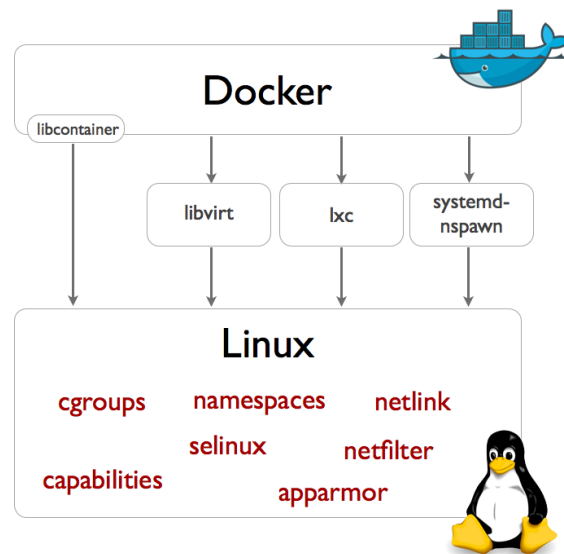
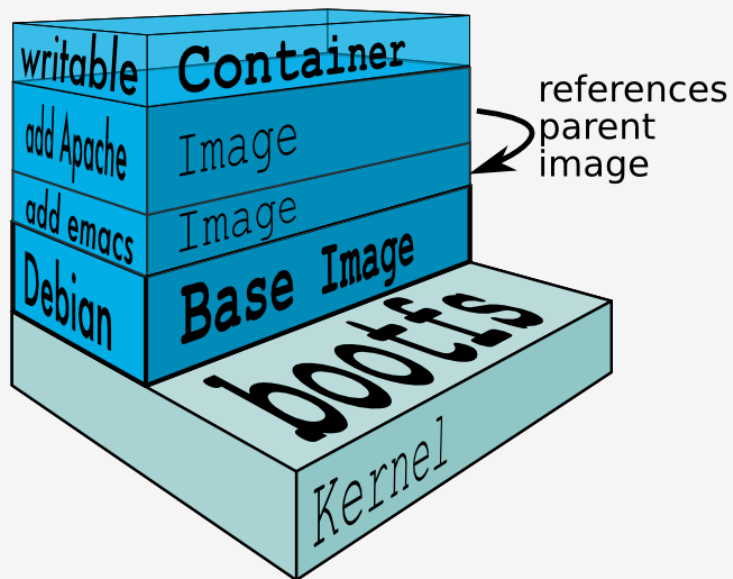
Network  
Isolation

Change  
Management

Sharing

Process  
Management

# Docker Components



# Docker Grounds up: Resource Isolation

## Cgroups : Isolation and accounting

- cpu
- memory
- block i/o
- devices
- network
- numa
- freezer

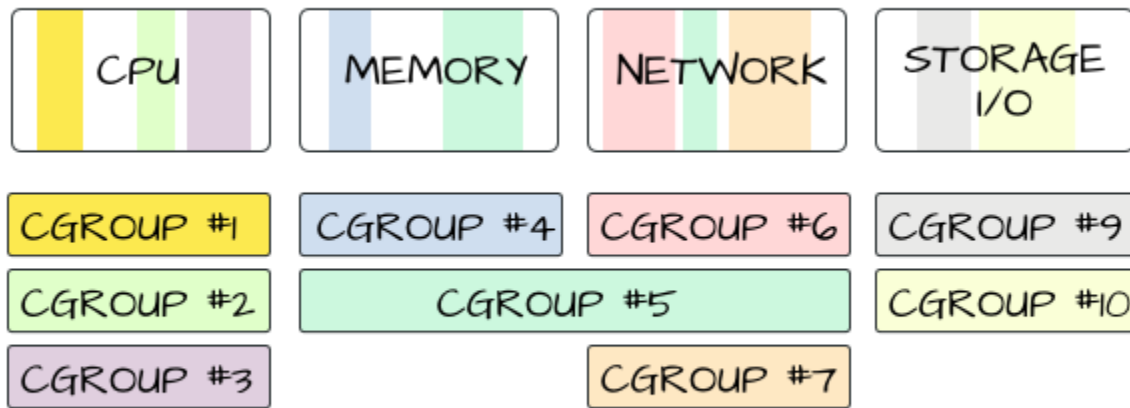


image credit: [mairin](#)

# Docker Grounds up: Namespaces

- Process trees.
- Mounts.
- Network.
- User accounts.
- Hostnames.
- Inter-process communication.

```
pid_t pid = clone(..., flags, ...)
```

CLONE_NEWUTS	hostname,
domainname	
CLONE_NEWIPC	IPC objects
CLONE_NEWPID	Process IDs
CLONE_NEWNET	Network
configuration	
CLONE_NEWNS	File system mounts
CLONE_NEWUSER	User and
Group IDs	

```
setns(int fd, int nstype)
```

```
CLONE_NEWIPC  
CLONE_NEWNET  
CLONE_NEWUTS
```

Also: `unshare(flags)`



# Docker Grounds up: Add Security

- Linux Capabilities
  - Drops most capabilities.
  - Enable what a task needs.
- GRSEC and PAX
- SELinux
- AppArmor

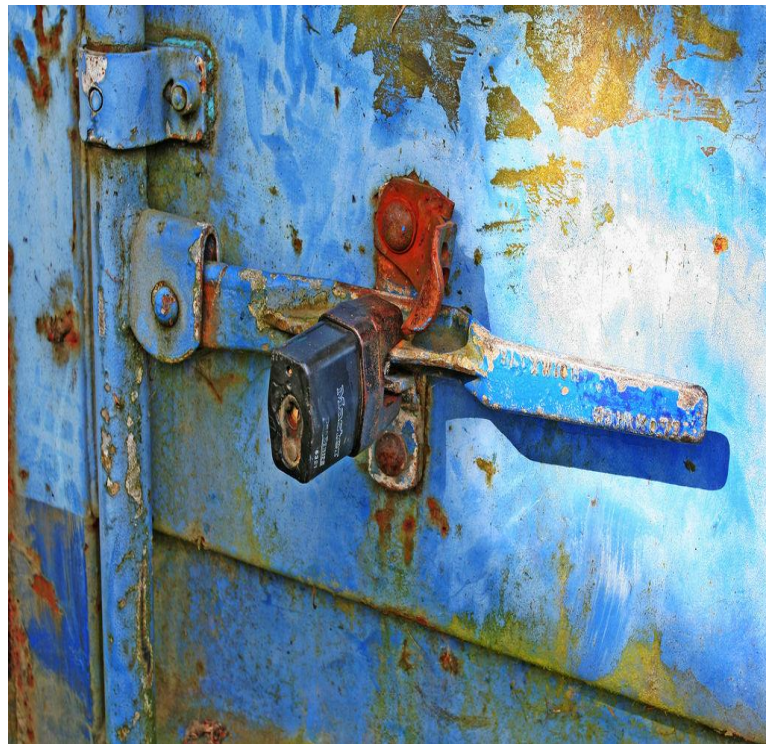


image credit: [Leo Reynolds](#)



# Docker Grounds up: Filesystem

## File-system Isolation:

Building a rootfs dir and chroot into it.

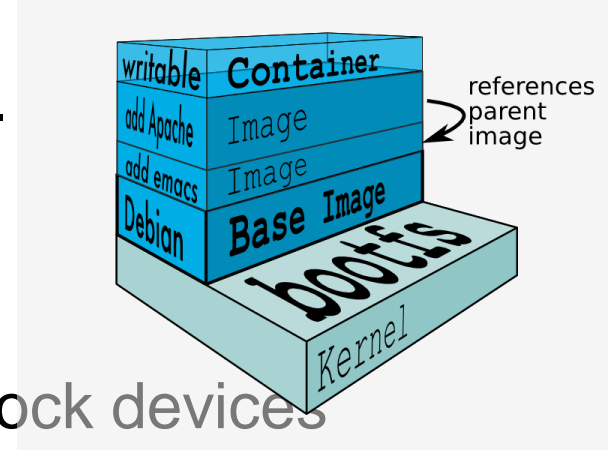
With mount namespace, use pivot-root.

## Features:

Layering, CoW, Caching, Diffing

## Solutions:

UnionFS, Snapshotting FS, CoW block devices



# Docker Grounds up: Filesystem

	Union Filesystems	Snapshotting Filesystems	Copy-on-write block devices
Provisioning	Superfast Supercheap	Fast Cheap	Fast Cheap
Changing small files	Superfast Supercheap	Fast Cheap	Fast Costly
Changing large files	Slow (first time) Inefficient (copy-up!)	Fast Cheap	Fast Cheap
Diffing	Superfast	Superfast	Slow
Memory usage	Efficient	Efficient	Inefficient (at high densities)
Drawbacks	Random quirks AUFS not mainline !AUFS more quirks	ZFS not mainline BTRFS not as nice	Higher disk usage Great performance (except diffing)
Bottom line	Ideal for PAAS and high density things	This is the Future (probably)	Dodge Ram 3500

From: Jérôme Petazzoni

# Docker Grounds up: Processes & Networking

We have resources, isolation, and file system management.

Docker daemon handles starting/stopping processes with:

- Attach logic

- Logs

- TTY management

- Docker run options

- Events and container state

Network Management

- NAT, Bridge, Veth

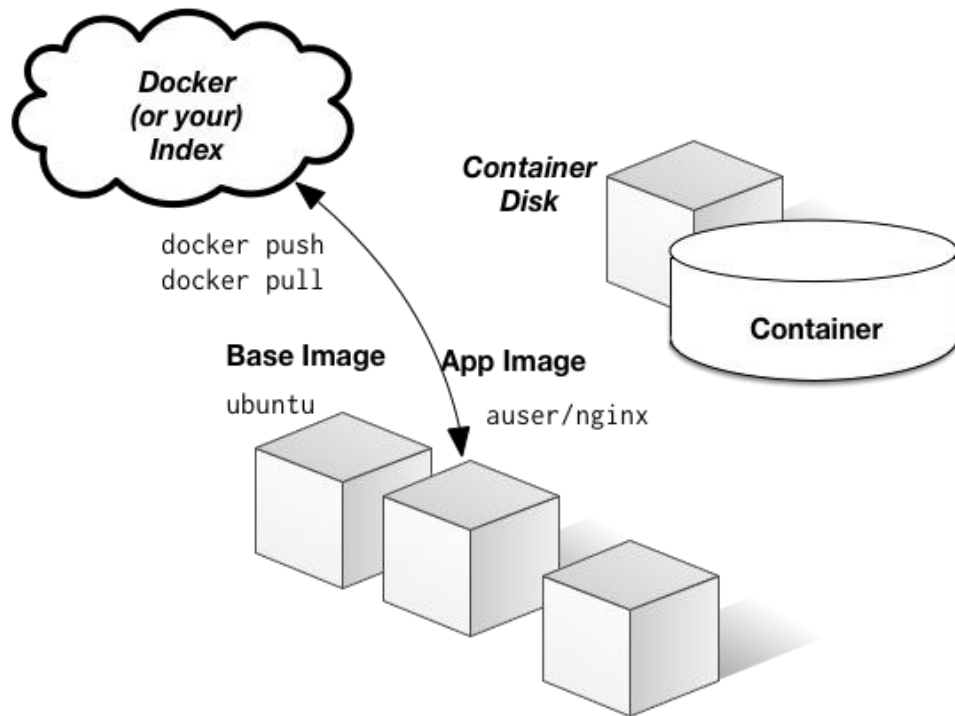
- Expose

- Links

# Docker Grounds up: Images

Create and share images  
Push, pull, commit images.  
Registry (public, private) and index.  
Dockerfiles

Orchestration:  
Linking Containers  
Multi-host linking  
Dynamic discovery



[image: jbarratt](#)

# Docker Codewalk

[github.com/dotcloud/docker/](https://github.com/dotcloud/docker/)

api : docker client and server api

daemon : Managing containers and images

engine: commands/jobs processing

graph: store for versioned filesystem images and their relationship.

registry: handling registry and repository.

links: Linking containers.

**integration-cli**: Integration tests.

**docs**: documentation.

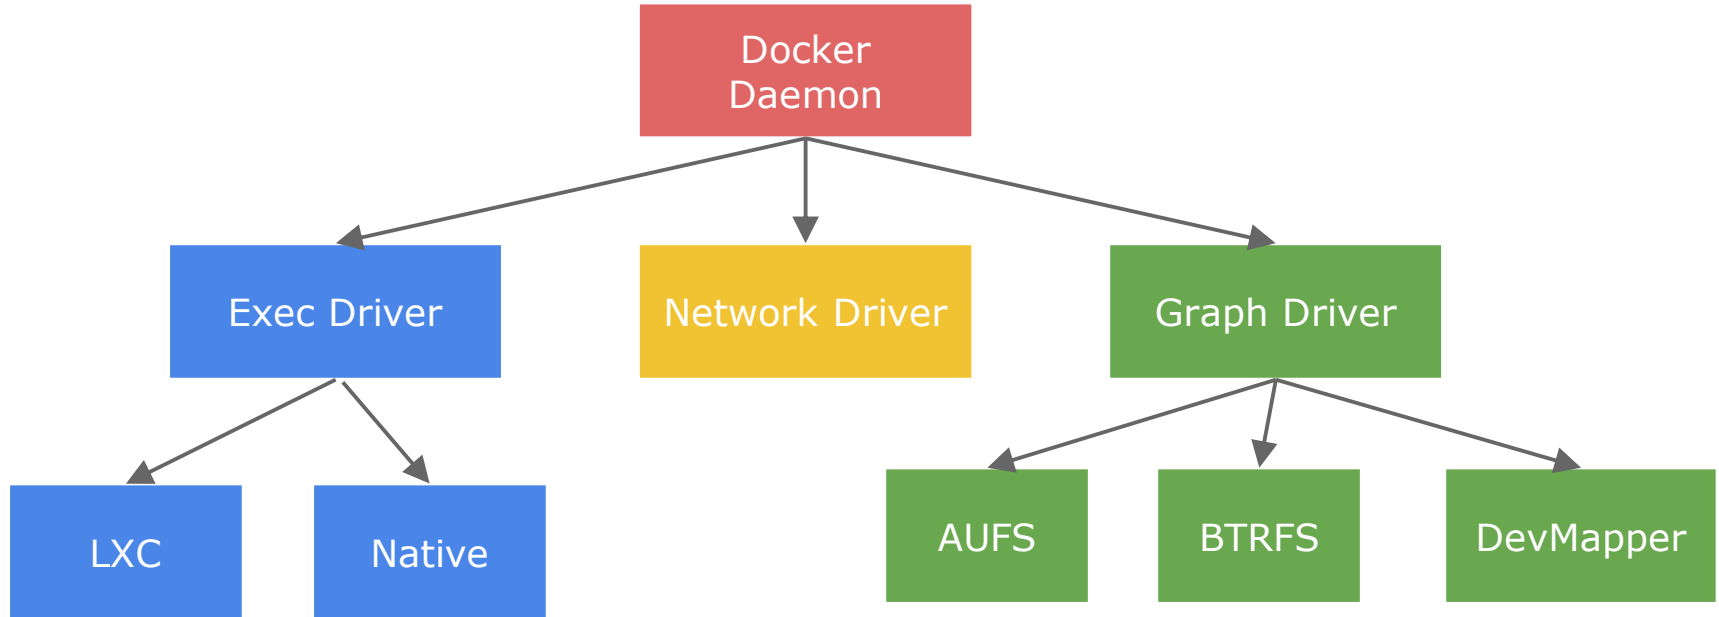
**pkg**: collection of standalone utility packages that are not docker specific.

-> Great place to start contributing.

Time for actual walkthrough...



# Docker Codewalk : docker/daemon



# Docker Codewalk : pkg

[github.com/dotcloud/docker/pkg](https://github.com/dotcloud/docker/pkg)

libcontainer: cgroup and namespaces. Uses lot of other utility packages.  
nsinit binary.

apparmor, selinux, label : applying security profiles.

mount, signals : system utilities.

iptables, networkfs, netlink : network utilities.

term: terminal handling

systemd

Let's look through some of these.

Thanks!

Rohit Jnagal

jnagal@google  
@jnagal



## Kesden Additional Slide:

- VMs vs Containers
  - VMs virtualize Hardware, OS, etc
  - Containers virtualize application environment
- Containers may not provide as strong a security model
  - What is virtualized? What is real?
  - What about the super users?
  - Things that are sideways, e.g. virtual file system, devices, etc
- Generally use containers to virtualize for one application in shared host or VM
- Use VMs to virtualize for many applications
- VMs probably 2-3x as resource intensive as containers
  - Corollary: Can get 2-3x as much from containerized solution vs VMs